Amazon Lightsail instance health metrics

Note
Lightsail provides several health metrics about your virtual private server (or Lightsail *instance*) to help you keep your application running smoothly. You can quickly look at CPU utilization, data transfer rates, and status check metrics over a specified time period.

Last updated: March 22, 2018

You can view the following metrics over different time periods.

CPUUtilization
This metric refers to the percentage of CPU capacity that your Lightsail instance is using. This metric identifies the processing power required to run an application on a selected instance.
If you experience a high degree of CPU utilization, your instance might be unable to handle the load that you're asking it to carry. Maybe your blog hit the big time? To fix a high-percentage value on **CPUUtilization**, you can add more CPUs or use a Lightsail instance with a faster CPU. You can also get more information by using the `top` command on Linux. It shows you all the processes running on your instance and the CPU and memory resources they're using.

NetworkIn
The number of bytes received on all network interfaces by the Lightsail instance. This metric identifies the volume of incoming network traffic to an application on a single instance.
With **NetworkOut**, this metric helps you measure data transfer rates. Data transfer rates are important when it comes to billing. For more information about your bill, see Lightsail Billing.
For troubleshooting, here are some Linux commands you can use to run network input and output diagnostics on your instance: 18 commands to monitor network bandwidth on Linux server.

NetworkOut
The number of bytes sent out on all network interfaces by the Lightsail instance. This metric identifies the volume of outgoing network traffic to an application on a single instance. With **NetworkIn**, this metric helps you measure data transfer rates.
For troubleshooting, here are some Linux commands you can use to run network input and output diagnostics on your instance: 18 commands to monitor network bandwidth on Linux server.

StatusCheckFailed
A combination of **StatusCheckFailed_Instance** and **StatusCheckFailed_System** that reports if either of the status checks has failed. Values for this metric are either zero (0) or one (1). A zero indicates that the status check passed. A one indicates a status check failure.
Status check metrics are available at a one-minute frequency. For a newly launched instance, status check metric data is available only after the instance completes the **Pending** state. Status check metrics become available within a few minutes of the instance being in the running state.

StatusCheckFailed_Instance
Reports whether the instance passed the Lightsail instance status check in the last minute. Values for this metric are either zero (0) or one (1). A zero indicates that the status check passed. A one indicates a status check failure.

Status check metrics are available at a one-minute frequency. For a newly launched instance, status check metric data is available only after the instance completes the **Pending** state. Status check metrics become available within a few minutes of the instance being in the running state.

StatusCheckFailed_System
Reports whether the instance passed the Lightsail system status check in the last minute. Values for this metric are either zero (0) or one (1). A zero indicates that the status check passed. A one indicates a status check failure.
Status check metrics are available at a one-minute frequency. For a newly launched instance, status check metric data is available only after the instance completes the **Pending** state. Status check metrics become available within a few minutes of the instance being in the running state.

For more information, see Amazon EC2 Metrics and Dimensions.